

# LUNAR LANDER

Saurabh Kumar SC22B146

Sahil Shete SC22B113

Satyajeet Musmade SC22B182

## REWARD MAXIMIZERS




INDIAN INSTITUTE OF SPACE SCIENCE AND TECHNOLOGY  
THIRUVANANTHAPURAM, KERALA, INDIA 695547

December 16, 2025



# Outline



## 1 Introduction

## 2 Model and Problem Statement

## 3 Reinforcement Learning Agent

## 4 Extensions

## 5 Performance Analysis



# Motivation for the Problem

- **Control Theory Challenge:** Landing a spacecraft is a classic problem requiring precision control under dynamic forces (gravity, thrust).
- **High-Dimensional Control:** Traditional methods become complex as system dynamics increase. Reinforcement Learning (RL) offers a model-free approach.
- **Deep Reinforcement Learning (DRL):** Merges the perception capability of Deep Learning with the decision-making of RL.
- **LunarLander-v3:** A canonical, yet challenging, benchmark environment for testing DRL algorithms like DQN.

## The Problem

Train an agent to safely navigate and land a lunar module on a designated landing spot, maximizing cumulative reward while minimizing fuel usage.



# Environment Description: LunarLander-v3

The problem is modeled as a Markov Decision Process (MDP)  
 $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$ .

- **State Space ( $\mathcal{S}$ )**: 8 continuous values

$$\text{Box}([ -2.5, -2.5, -10, -10, -6.2831855, -10, 0, 0 ], [ 2.5, 2.5, 10, 10, 6.2831855, 10, 1, 1 ])$$


i.e.,  $x$  position,  $y$  position,  $x$  velocity,  $y$  velocity, angle, angular velocity, left leg contact, right leg contact.

- **Action Space ( $\mathcal{A}$ )**: 4 discrete actions

- 1 Do nothing
- 2 Fire main engine
- 3 Fire left engine
- 4 Fire right engine

## Performance Criterion (Goal)

Achieve an average score of  $\geq 200.0$  over 100 consecutive episodes.



**Figure 1:** The LunarLander-v3 gym environment





# Performance Criterion: Reward Structure

- **Landing/Success:** +100 points for coming to rest on the landing pad.
- **Crashing/Failure:** -100 points for contacting the ground outside the pad or non-zero speed.
- **Fuel Usage (Cost):** -0.3 points for firing the main engine, -0.03 for side engines.
- **Leg Contact:** +10 points for each leg in contact with the ground.
- **Discount Factor:** 0.99.

## The Return

The agent maximizes the discounted cumulative reward (Return):


$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$



# The Deep Q-Network (DQN) Agent

## Agent Architecture:

- Uses two identical **Feedforward Neural Networks** (Local and Target)
- Layers:  
Input(8) → FC(64, ReLU) → FC(64, ReLU) → Output(4)
- **Experience Replay:** Stores experiences in a capacity =  $10^5$  buffer
- **$\epsilon$ -Greedy Policy:** Balances exploration/exploitation



**Figure 2: DQN Architecture (Local and Target Networks)**



# DQN Learning and Updates

- **Q-Target Calculation (Bellman Equation):**

$$Y_t = R_{t+1} + \gamma \max_a Q_{\text{target}}(S_{t+1}, a_{t+1}) \cdot (1 - \text{Done})$$


- **Loss Function:** Mean Squared Error (MSE) between the target and the local network's estimate.

$$\mathcal{L} = \mathbb{E}[(Y_t - Q_{\text{local}}(S_t, A_t))^2]$$

- **Soft Update of Target Network:** Stabilizes training by slowly updating the target weights ( $\theta_{\text{target}}$ ).

$$\theta_{\text{target}} \leftarrow \tau \theta_{\text{local}} + (1 - \tau) \theta_{\text{target}}$$

- $\tau$  (interpolation\_parameter) =  $10^{-3}$ .
- Updates occur every 4 steps after sampling a minibatch (size = 100).



**Figure 3:** DQN Training process




# Key Hyperparameters

- **Learning Rate (Adam Optimizer):**  $5 \times 10^{-4}$
- **Discount Factor:** 0.99
- **Replay Buffer Size:**  $10^5$
- **$\epsilon$ -Decay Schedule:**
  - Starting  $\epsilon$ : 1.0
  - Ending  $\epsilon$ : 0.01
  - Decay Rate: 0.995




# Performance Analysis



**Figure 4:** Episode Reward vs. Episode No.

# Performance Analysis



**Figure 5:** Training Loss vs. Episode No.



Thank you

---