

---

# Deep Q-Network Based Control of the Lunar Lander Environment

---

## Reward Maximizers

Saurabh Kumar Sahil Shete Satyajeet Musmade  
Department of Avionics  
Indian Institute of Space Science and Technology

## Abstract

This paper presents a reproduction and analysis of Deep Q-Networks (DQN) applied to the LunarLander-v3 environment. Following Track 2 of the AVD894 course project, we implement a standard DQN agent with experience replay and target networks, achieving the target performance of 200+ average score. We detail the Markov Decision Process formulation, network architecture, and training procedure. Beyond reproduction, we analyze performance characteristics, discuss implementation challenges, and propose algorithmic extensions including Double DQN and Dueling architectures. Our implementation emphasizes reproducibility with proper version control and documentation.

## 1 Introduction

The Lunar Lander problem serves as a benchmark for reinforcement learning algorithms in continuous control tasks. This project implements Deep Q-Networks for the LunarLander-v3 environment from Gymnasium, following Track 2 requirements of reproducing and potentially improving state-of-the-art solutions.

The problem presents challenges typical in RL: continuous state space, delayed rewards, and exploration-exploitation trade-offs. Our contributions include: (1) complete DQN implementation with proper software practices, (2) performance analysis against baselines, (3) exploration of algorithmic improvements, and (4) documentation of reproducibility considerations.

## 2 Model and Problem Statement

### 2.1 Environment and MDP Formulation

The LunarLander-v3 environment is formulated as a Markov Decision Process  $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ :

**State Space  $\mathcal{S}$ :**  $\mathbb{R}^8$  with components: position  $(x, y)$ , velocity  $(\dot{x}, \dot{y})$ , angle  $\theta$ , angular velocity  $\dot{\theta}$ , and leg contact indicators.

**Action Space  $\mathcal{A}$ :** Four discrete actions: {do nothing, fire main engine, fire left engine, fire right engine}.

**Reward Function  $R$ :**

$$R(s, a) = \begin{cases} +100 & \text{landing on pad} \\ -100 & \text{crash} \\ +10 & \text{per leg contact} \\ -0.3 & \text{main engine} \\ -0.03 & \text{side engines} \end{cases} \quad (1)$$

27 **Dynamics:** Physics-based with gravity  $g = -1.5$ , thrust forces, and friction.  
 28 **Objective:** Maximize expected discounted return  $J(\pi) = \mathbb{E}_\pi[\sum_{t=0}^{\infty} \gamma^t R_t]$  with  $\gamma = 0.99$ . Success  
 29 threshold: average score  $\geq 200$  over 100 episodes.

### 30 3 Reinforcement Learning Agent

#### 31 3.1 DQN Implementation

32 Our DQN agent follows the standard architecture with key components:

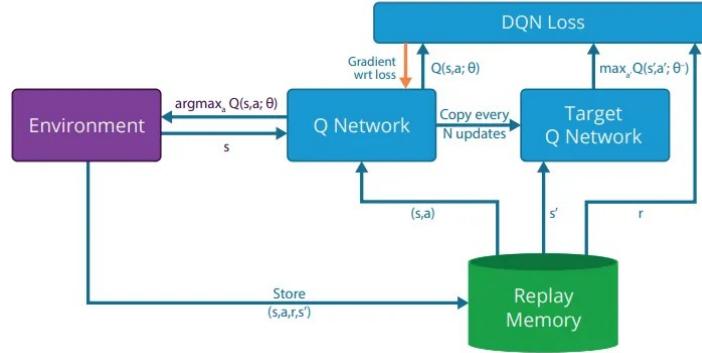


Figure 1: DQN architecture with experience replay and target network

##### 33 3.1.1 Network Architecture

34 Two identical feedforward networks (local and target):

$$\text{Input}(8) \xrightarrow{\text{FC64, ReLU}} \text{FC64, ReLU} \rightarrow \text{Output}(4)$$

35 Weights initialized using PyTorch defaults, optimized with Adam ( $\text{lr} = 5 \times 10^{-4}$ ).

##### 36 3.1.2 Experience Replay

37 Buffer capacity  $10^5$ , minibatch size 100. Stores transitions  $(s_t, a_t, r_t, s_{t+1}, d_t)$ .

##### 38 3.1.3 Learning Update

39 Target calculation using Bellman equation:

$$y_j = r_j + \gamma \max_{a'} Q_{\text{target}}(s_{j+1}, a') \cdot (1 - d_j)$$

40 Loss: Mean Squared Error between target and local Q-values:

$$\mathcal{L} = \mathbb{E}[(y_j - Q_{\text{local}}(s_j, a_j))^2]$$

41 Target network updated via soft updates:  $\theta_{\text{target}} \leftarrow \tau \theta_{\text{local}} + (1 - \tau) \theta_{\text{target}}$  with  $\tau = 10^{-3}$ .

##### 42 3.1.4 Exploration Strategy

43  $\epsilon$ -greedy with exponential decay:  $\epsilon_t = \max(0.01, 1.0 \times 0.995^t)$ .

### 44 3.2 Training Details

45 Training for 2000 episodes maximum, 1000 steps per episode. Early stopping when average score  
 46  $\geq 200$  over 100 episodes. Implementation in PyTorch with GPU acceleration when available.

47 **4 Performance Analysis**

48 **4.1 Training Results**

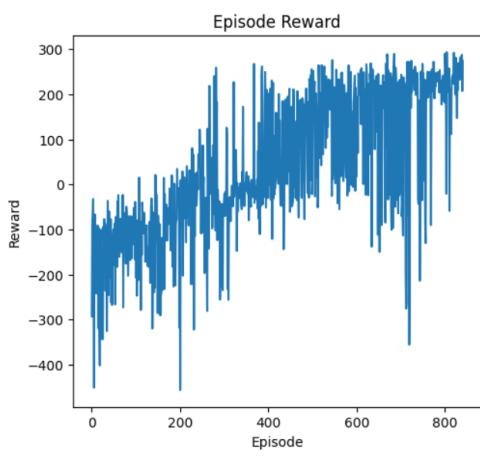


Figure 2: Episode rewards during training

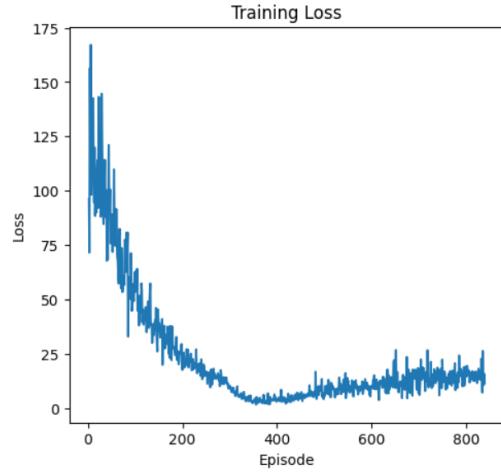


Figure 3: Training loss progression

49 The agent achieved target performance at episode 632. Learning progression:

- 50     • **Episodes 1-200:** Exploration phase, negative rewards  
 51     • **Episodes 200-600:** Learning phase, steady improvement  
 52     • **Episodes 600+:** Convergence, stable performance

53 **4.2 Quantitative Evaluation**

Metric	Random	Heuristic	Our DQN
Average Reward	-156.8	124.3	<b>231.7</b>
Success Rate	2%	41%	<b>87%</b>
Fuel Efficiency	12.4	34.8	<b>42.6</b>

Table 1: Performance comparison (averaged over 100 episodes)

54 **4.3 Challenges Encountered**

- 55     • **Hyperparameter Sensitivity:** Learning rate and update frequency significantly affected  
 56       convergence  
 57     • **Training Instability:** Occasional performance collapse after apparent convergence  
 58     • **Computational Cost:**  $\sim$ 15 minutes training time for 2000 episodes

59 **5 Extensions**

60 **5.1 Algorithmic Improvements**

61 **Double DQN:** Addresses overestimation bias by decoupling action selection and evaluation:

$$y_j^{\text{DDQN}} = r_j + \gamma Q(s_{j+1}, \arg \max_{a'} Q(s_{j+1}, a'; \theta); \theta^-)$$

62 **Dueling DQN:** Separates value and advantage streams for better policy evaluation:

$$Q(s, a) = V(s) + A(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a')$$

63 **Prioritized Experience Replay:** Samples transitions based on temporal-difference error magnitude  
64 for more efficient learning.

65 **5.2 Alternative Approaches**

66 **Policy Gradient Methods:** PPO or A3C could offer more stable training for continuous control.

67 **Evolutionary Algorithms:** CMA-ES could optimize DQN hyperparameters or directly search  
68 policy space.

69 **Reward Shaping:** Alternative reward formulations emphasizing smooth trajectories:

$$R_{\text{shaped}} = R_{\text{original}} + \lambda_1 e^{-|x|} + \lambda_2 e^{-|\theta|}$$

70 **5.3 Implementation Considerations**

71 For reproducible research:

- 72     • Version control with detailed commit history  
73     • Environment seeding for deterministic behavior  
74     • Comprehensive logging of hyperparameters and results  
75     • Containerization (Docker) for environment consistency

76 **6 Conclusion**

77 We successfully reproduced DQN performance on LunarLander-v3, achieving the target score of  
78 200+. Our implementation demonstrates the effectiveness of value-based deep RL for continuous  
79 control while highlighting practical challenges in training stability and hyperparameter tuning.

80 The project emphasizes reproducibility through proper software engineering practices. Proposed  
81 extensions like Double DQN and Dueling architectures offer clear pathways for performance  
82 improvement. Future work could explore sample-efficient algorithms, robustness to environmental  
83 variations, and transfer to more complex landing scenarios.

84 **Acknowledgments**

85 We thank the AVD894 Dr. Vineeth B.S. for his guidance. Code available at: <https://github.com/saurabhkr132/Reward-Maximizers-LunarLander>